

- 0 On all three machines, install Docker with the command: `sudo yum install docker -y`
- 0 Configure Docker to use `systemd` as the `cgroup` driver by creating and editing the `daemon.json` file:
  - Change directory to `/etc/docker`

- Use the command `cat <<EOF | sudo tee /etc/docker/daemon.json` followed by the JSON configuration details and end with EOF

- Enable and restart Docker: `sudo systemctl enable docker sudo systemctl daemon-reload sudo systemctl restart docker docker -v`

```
[ec2-user@ip-172-31-92-18 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:09:56 ago on Wed Sep 11 15:19:39 2024.
Dependencies resolved.
```

Package	Architecture
Installing:	
docker	x86_64
Installing dependencies:	
containerd	x86_64
iptables-libs	x86_64
iptables-nft	x86_64
libcgroup	x86_64
libnetfilter_conntrack	x86_64
libnfnetlink	x86_64
libnftnl	x86_64
pigz	x86_64
runc	x86_64

Transaction Summary

#### 4. Kubernetes Installation:

- Disable SELinux before configuring kubelet: `sudo setenforce 0 sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config`
- Add the Kubernetes repository and install Kubernetes components:
  - Use the command `cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo` followed by the repository configuration details and end with EOF `sudo yum update sudo yum install -y kubelet kubeadm kubectl 0--disableexcludes=kubernetes`

- Configure networking for bridging: `sudo swapoff -a` `echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf` `sudo sysctl -p`

```
[ec2-user@ip-172-31-81-63 docker]$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
docker -v

Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
Docker version 25.0.5, build 5dc9bcc
[ec2-user@ip-172-31-81-63 docker]$
```

```
[ec2-user@ip-172-31-81-63 docker]$ sudo setenforce 0
[ec2-user@ip-172-31-81-63 docker]$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

```
[ec2-user@ip-172-31-81-63 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Last metadata expiration check: 0:01:34 ago on Wed Sep 11 15:39:05 2024.
Dependencies resolved.
=====
Package                                Architecture      Version
=====
Installing:
kubeadm                                x86_64            1.30.4-150500.1.1
kubectl                                x86_64            1.30.4-150500.1.1
kubelet                                x86_64            1.30.4-150500.1.1
Installing dependencies:
conntrack-tools                        x86_64            1.4.6-2.amzn2023.0.2
cri-tools                              x86_64            1.30.1-150500.1.1
kubernetes-cni                         x86_64            1.4.0-150500.1.1
libnetfilter_cthelper                  x86_64            1.0.0-21.amzn2023.0.2
libnetfilter_cttimeout                 x86_64            1.0.0-19.amzn2023.0.2
libnetfilter_queue                     x86_64            1.0.5-2.amzn2023.0.2
socat                                   x86_64            1.7.4.2-1.amzn2023.0.2
=====
Transaction Summary
=====
Install 10 Packages
```

## 5. Master Node Setup:

- Initialize the Kubernetes master node (perform only on the master machine): `sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all`
- After initialization, set up the Kubernetes configuration on the master node: `mkdir -p $HOME/.kube` `sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/` `sudo chown $(id -u):$(id -g) $HOME/.kube/` `config`

- Save the generated join command from the output for worker nodes. This command is unique and specific to your cluster setup: `kubeadm join 172.31.91.120:6443 --token r8j60r.nlj6h0klbewvoka5\--discovery-token-ca-cert-hashsha256:dd8426260174d673303aef17717f740772fcf7ee782245bc653eecf4a13 05da7`

```
ubuntu@ip-172-31-28-117:~$ sudo kubeadm join 172.31.27.176:6443 --token ttay2x.n0squeukjai8sgfg3 \
--discovery-token-ca-cert-hash sha256:d6fc5fb7e984c83e2807780047fec6c4f2acfe9da9184ecc028d77157608fbb6

[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 501.396793ms
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

- Deploy the Flannel networking plugin to enable pod communication: `kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml`
- Check the status of the pods to ensure they are running: `kubectl get pods --all-namespaces`

```
ubuntu@ip-172-31-27-176:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-18-135	NotReady	<none>	88s	v1.31.1
ip-172-31-27-176	NotReady	control-plane	10m	v1.31.1
ip-172-31-28-117	NotReady	<none>	2m58s	v1.31.1

```
ubuntu@ip-172-31-27-176:~$ kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
```

## sudo systemctl status kubelet

```
ubuntu@ip-172-31-27-176:~$ sudo systemctl status kubelet
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/usr/lib/systemd/system/kubelet.service; enabled; preset: enabled)
   Drop-In: /usr/lib/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Mon 2024-09-16 15:40:01 UTC; 11min ago
     Docs: https://kubernetes.io/docs/
   Main PID: 5989 (kubelet)
    Tasks: 10 (limit: 4676)
   Memory: 32.6M (peak: 33.2M)
      CPU: 10.705s
   CGroup: /system.slice/kubelet.service
            └─5989 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --config=/var/

Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497458      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497516      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497569      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497620      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497669      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497719      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s
Sep 16 15:51:31 ip-172-31-27-176 kubelet[5989]: E0916 15:51:31.605091      5989 kubelet.go:2902] "Container runtime network not ready" networkReady="NetworkR
Sep 16 15:51:32 ip-172-31-27-176 kubelet[5989]: I0916 15:51:32.366237      5989 scope.go:117] "RemoveContainer" containerID="f44f06967c5b3e567e07841a7b4352ae"
Sep 16 15:51:36 ip-172-31-27-176 kubelet[5989]: E0916 15:51:36.606675      5989 kubelet.go:2902] "Container runtime network not ready" networkReady="NetworkR
Sep 16 15:51:41 ip-172-31-27-176 kubelet[5989]: E0916 15:51:41.608404      5989 kubelet.go:2902] "Container runtime network not ready" networkReady="NetworkR
```

Now Run command `kubectl get nodes -o wide` we can see Status is ready.

```
ubuntu@ip-172-31-27-176:~$ kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME
ip-172-31-18-135	Ready	<none>	6m19s	v1.31.1	172.31.18.135	<none>	Ubuntu 24.04 LTS	6.8.0-1012-aws	containerd://1.7.12
ip-172-31-27-176	Ready	control-plane	15m	v1.31.1	172.31.27.176	<none>	Ubuntu 24.04 LTS	6.8.0-1012-aws	containerd://1.7.12
ip-172-31-28-117	Ready	<none>	7m49s	v1.31.1	172.31.28.117	<none>	Ubuntu 24.04 LTS	6.8.0-1012-aws	containerd://1.7.12

## 6. Worker Node Setup:

- On each worker node, install the required package and configure kubelet: `sudo yum install iproute-tc -y sudo systemctl enable kubelet sudo systemctl restart kubelet`
- Join the worker nodes to the Kubernetes cluster using the join command from the master node: `kubeadm join 172.31.91.120:6443 --token r8j60r.n1j6h0klbewvoka5\--discovery-token-ca-cert-hashsha256:dd8426260174d673303aef17717f740772fcf7ee782245bc653eecf4a1305da7`

## Node 1

```
ubuntu@ip-172-31-28-117:~$ sudo kubeadm join 172.31.27.176:6443 --token ttay2x.n0squeukjai8sgfg3 \
--discovery-token-ca-cert-hash sha256:d6fc5fb7e984c83e2807780047fec6c4f2acfe9da9184ecc028d77157608fbb6

[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 501.396793ms
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@ip-172-31-28-117:~$
```

## Node 2

```
ubuntu@ip-172-31-18-135:~$ sudo kubeadm join 172.31.27.176:6443 --token ttay2x.n0squeukjai8sgfg3 \
--discovery-token-ca-cert-hash sha256:d6fc5fb7e984c83e2807780047fec6c4f2acfe9da9184ecc028d77157608fbb6

[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.001003808s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@ip-172-31-18-135:~$
```

## 7. Verify Node Status:

- On the master node, verify that the worker nodes have successfully joined the cluster by running: `watch kubectl get nodes`

Or run `kubectl get nodes`

```
ubuntu@ip-172-31-27-176:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
ip-172-31-18-135    Ready    Node2    24m   v1.31.1
ip-172-31-27-176    Ready    control-plane  33m   v1.31.1
ip-172-31-28-117    Ready    Node1    25m   v1.31.1
ubuntu@ip-172-31-27-176:~$
```

## Conclusion:

Setting up the Kubernetes cluster involved several challenges. Network configuration issues initially hindered the deployment of the Flannel plugin, requiring open ports and a functional Kubernetes API server. Disabling SELinux and adjusting firewall rules were essential for proper communication between

components. Worker nodes experienced difficulties with the kubelet service, which needed to be correctly configured and restarted. Additionally, accurate copying of the join command, including the token and discovery-token-ca-cert-hash, was crucial for integrating worker nodes into the cluster. These issues underscored the need for precise configuration and troubleshooting to achieve a stable Kubernetes setup.